



Algorithm Design V

Divide and Conquer II

Guoqiang Li
School of Computer Science



SHANGHAI JIAO TONG
UNIVERSITY

Reviews of the Previous Lecture



The **divide-and-conquer** strategy solves a problem by:

- 1 Breaking it into **subproblems** that are themselves smaller instances of the same type of problem.
- 2 **Recursively** solving these subproblems.
- 3 Appropriately **combining** their answers.



Master Theorem

If $T(n) = aT(\lceil n/b \rceil) + O(n^d)$ for some constants $a > 0$, $b > 1$ and $d \geq 0$, then

$$T(n) = \begin{cases} O(n^d) & \text{if } d > \log_b a \\ O(n^d \log n) & \text{if } d = \log_b a \\ O(n^{\log_b a}) & \text{if } d < \log_b a \end{cases}$$

An Exercise



Let B be an $n \times n$ chessboard, where n is a power of 2. Use a divide-and-conquer argument to describe how to cover all squares of B except one with L -shaped tiles. For example, if $n = 2$, then there are four squares three of which can be covered by one L -shaped tile, and if $n = 4$, then there are 16 squares of which 15 can be covered by 5 L -shaped tiles.

An Exercise



- I. 如果图 G 小于 6 个节点, 直接暴力枚举求出其最小割。 $\leftarrow$
- II. 采用随机收缩边的方式将原图的节点收缩至 $n/\sqrt{2}$ 个节点得到图 G_1 。 $\leftarrow$
- III. 采用随机收缩边的方式将原图的节点收缩至 $n/\sqrt{2}$ 个节点得到图 G_2 。 $\leftarrow$
- IV. 对图 G_1 和 G_2 分别调用该算法并取其中更好的结果。 $\leftarrow$

其中比较时间需要的时间复杂度是 $O(n^2)$ ， n 是图 G 的节点数目。试问，整个问题的复杂度是多少？^①

Counting Inversions

Counting Inversions

Music site tries to match your song preferences with others.

- You rank n songs.
- Music site consults database to find people with similar tastes.

Similarity metric: number of **inversions** between two rankings.

- My rank: $1, 2, \dots, n$.
- Your rank: $a_1, a_2, \dots, a_n$.
- Songs i and j are inverted if $i < j$, but $a_i > a_j$.

	A	B	C	D	E
me	1	2	3	4	5
you	1	3	4	2	5

2 inversions: 3-2, 4-2

Divide-and-Conquer



SHANGHAI JIAO TONG
UNIVERSITY

Divide-and-Conquer



Divide: separate list into two halves A and B .

Divide-and-Conquer



Divide: separate list into two halves A and B .

Conquer: recursively count inversions in each list.

Divide-and-Conquer



Divide: separate list into two halves A and B .

Conquer: recursively count inversions in each list.

Combine: count inversions (a, b) with $a \in A$ and $b \in B$.

Divide-and-Conquer



Divide: separate list into two halves A and B .

Conquer: recursively count inversions in each list.

Combine: count inversions (a, b) with $a \in A$ and $b \in B$.

Return sum of three counts.



Divide-and-Conquer

Divide: separate list into two halves A and B .

Conquer: recursively count inversions in each list.

Combine: count inversions (a, b) with $a \in A$ and $b \in B$.

Return sum of three counts.

input

1	5	4	8	10	2	6	9	3	7
---	---	---	---	----	---	---	---	---	---



Divide-and-Conquer

Divide: separate list into two halves A and B .

Conquer: recursively count inversions in each list.

Combine: count inversions (a, b) with $a \in A$ and $b \in B$.

Return sum of three counts.

input

1	5	4	8	10	2	6	9	3	7
---	---	---	---	----	---	---	---	---	---

count inversions in left half A

1	5	4	8	10
---	---	---	---	----

5-4

count inversions in right half B

2	6	9	3	7
---	---	---	---	---

6-3 9-3 9-7



Divide-and-Conquer

Divide: separate list into two halves A and B .

Conquer: recursively count inversions in each list.

Combine: count inversions (a, b) with $a \in A$ and $b \in B$.

Return sum of three counts.

input

1	5	4	8	10	2	6	9	3	7
---	---	---	---	----	---	---	---	---	---

count inversions in left half A

1	5	4	8	10
---	---	---	---	----

5-4

count inversions in right half B

2	6	9	3	7
---	---	---	---	---

6-3 9-3 9-7

count inversions (a, b) with $a \in A$ and $b \in B$

1	5	4	8	10
---	---	---	---	----

2	6	9	3	7
---	---	---	---	---

5-2 5-3 4-2 4-3 8-2 8-3 8-6 8-7 10-2 10-3 10-6 10-7 10-9



Divide-and-Conquer

Divide: separate list into two halves A and B .

Conquer: recursively count inversions in each list.

Combine: count inversions (a, b) with $a \in A$ and $b \in B$.

Return sum of three counts.

input

1	5	4	8	10	2	6	9	3	7
---	---	---	---	----	---	---	---	---	---

count inversions in left half A

1	5	4	8	10
---	---	---	---	----

5-4

count inversions in right half B

2	6	9	3	7
---	---	---	---	---

6-3 9-3 9-7

count inversions (a, b) with $a \in A$ and $b \in B$

1	5	4	8	10
---	---	---	---	----

2	6	9	3	7
---	---	---	---	---

5-2 5-3 4-2 4-3 8-2 8-3 8-6 8-7 10-2 10-3 10-6 10-7 10-9

output $1+3+13=17$

How to Combine Two Subproblems?



SHANGHAI JIAO TONG
UNIVERSITY

Q. How to count inversions (a, b) with $a \in A$ and $b \in B$?

How to Combine Two Subproblems?



Q. How to count inversions (a, b) with $a \in A$ and $b \in B$?

A. Easy if A and B are sorted!

How to Combine Two Subproblems?

Q. How to count inversions (a, b) with $a \in A$ and $b \in B$?

A. Easy if A and B are sorted!

Warmup algorithm.

- Sort A and B .
- For each element $b \in B$,
 - binary search in A to find how many elements in A are greater than b .



How to Combine Two Subproblems?

Q. How to count inversions (a, b) with $a \in A$ and $b \in B$?

A. Easy if A and B are sorted!

Warmup algorithm.

- Sort A and B .
- For each element $b \in B$,
 - binary search in A to find how many elements in A are greater than b .

list A

7	10	18	3	14
---	----	----	---	----

list B

20	23	2	11	16
----	----	---	----	----

How to Combine Two Subproblems?

Q. How to count inversions (a, b) with $a \in A$ and $b \in B$?

A. Easy if A and B are sorted!

Warmup algorithm.

- Sort A and B .
- For each element $b \in B$,
 - binary search in A to find how many elements in A are greater than b .

list *A*

7 10 18 3 14

sort A

3 7 10 14 18

list B

20 23 2 11 16

sort B

2 11 16 20 23



How to Combine Two Subproblems?

Q. How to count inversions (a, b) with $a \in A$ and $b \in B$?

A. Easy if A and B are sorted!

Warmup algorithm.

- Sort A and B .
- For each element $b \in B$,
 - binary search in A to find how many elements in A are greater than b .

list A

7	10	18	3	14
---	----	----	---	----

sort A

3	7	10	14	18
---	---	----	----	----

list B

20	23	2	11	16
----	----	---	----	----

sort B

2	11	16	20	23
---	----	----	----	----

binary search to count inversions (a, b) with $a \in A$ and $b \in B$

3	7	10	14	18
---	---	----	----	----

2	11	16	20	23
5	2	1	0	0

How to Combine Two Subproblems?



Count inversions (a, b) with $a \in A$ and $b \in B$, assuming A and B are sorted.

How to Combine Two Subproblems?



Count inversions (a, b) with $a \in A$ and $b \in B$, assuming A and B are sorted.

- Scan A and B from left to right.

How to Combine Two Subproblems?



Count inversions (a, b) with $a \in A$ and $b \in B$, assuming A and B are sorted.

- Scan A and B from left to right.
- Compare a_i and b_j .

How to Combine Two Subproblems?



Count inversions (a, b) with $a \in A$ and $b \in B$, assuming A and B are sorted.

- Scan A and B from left to right.
- Compare a_i and b_j .
- If $a_i < b_j$, then a_i is not inverted with any element left in B .

How to Combine Two Subproblems?



Count inversions (a, b) with $a \in A$ and $b \in B$, assuming A and B are sorted.

- Scan A and B from left to right.
- Compare a_i and b_j .
- If $a_i < b_j$, then a_i is not inverted with any element left in B .
- If $a_i > b_j$, then b_j is inverted with every element left in A .

How to Combine Two Subproblems?



Count inversions (a, b) with $a \in A$ and $b \in B$, assuming A and B are sorted.

- Scan A and B from left to right.
- Compare a_i and b_j .
- If $a_i < b_j$, then a_i is not inverted with any element left in B .
- If $a_i > b_j$, then b_j is inverted with every element left in A .
- Append smaller element to sorted list C .

How to Combine Two Subproblems?

Count inversions (a, b) with $a \in A$ and $b \in B$, assuming A and B are sorted.

- Scan A and B from left to right.
- Compare a_i and b_j .
- If $a_i < b_j$, then a_i is not inverted with any element left in B .
- If $a_i > b_j$, then b_j is inverted with every element left in A .
- Append smaller element to sorted list C .

count inversions (a, b) with $a \in A$ and $b \in B$

Diagram illustrating the mapping of elements from the first array to the second array. The first array contains elements 3, 7, 10, a_i , and 18. The second array contains elements 2, 11, b_j , 20, and 23. A blue arrow points from a_i to b_j , indicating the mapping.

How to Combine Two Subproblems?




Count inversions (a, b) with $a \in A$ and $b \in B$, assuming A and B are sorted.

- Scan A and B from left to right.
- Compare a_i and b_j .
- If $a_i < b_j$, then a_i is not inverted with any element left in B .
- If $a_i > b_j$, then b_j is inverted with every element left in A .
- Append smaller element to sorted list C .

count inversions (a, b) with $a \in A$ and $b \in B$



merge to form sorted list C



Algorithm Implementation



Sort-And-Count(L);

input : List L

output: Number of inversions in L and L in sorted order

if *List L has one element* **then**

 RETURN (0, L);

end

Divide the list into two halves A and B ;

$(r_A, A) \leftarrow \text{Sort-And-Count}(A)$;

$(r_B, B) \leftarrow \text{Sort-And-Count}(B)$;

$(r_{AB}, L) \leftarrow \text{Merge-And-Count}(A, B)$;

RETURN ($r_A + r_B + r_{AB}$, L);



Proposition

The sort-and-count algorithm counts the number of inversions in a permutation of size n in $O(n \log n)$ time.



Proposition

The sort-and-count algorithm counts the number of inversions in a permutation of size n in $O(n \log n)$ time.

Proof.



Proposition

The sort-and-count algorithm counts the number of inversions in a permutation of size n in $O(n \log n)$ time.

Proof.

$$T(n) = 2 \cdot T(\lceil n/2 \rceil) + \Theta(n)$$

Complex Number

Complex Number



$z = a + bi$ is plotted at position (a, b) .




Complex Number

$z = a + bi$ is plotted at position (a, b) .

In its polar coordinates, denoted (r, θ) , rewrite as

$$z = r(\cos \theta + i \sin \theta) = re^{i\theta}$$

- **length**: $r = \sqrt{a^2 + b^2}$.
- **angle**: $\theta \in [0, 2\pi)$.
- θ can always be reduced modulo 2π .






Complex Number

$z = a + bi$ is plotted at position (a, b) .

In its polar coordinates, denoted (r, θ) , rewrite as


$$z = r(\cos \theta + i \sin \theta) = re^{i\theta}$$

- **length**: $r = \sqrt{a^2 + b^2}$.
- **angle**: $\theta \in [0, 2\pi)$.
- θ can always be reduced modulo 2π .



Basic arithmetic:

- $-z = (r, \theta + \pi)$.
- $(r_1, \theta_1) \times (r_2, \theta_2) = (r_1 r_2, \theta_1 + \theta_2)$.
- If z is on the unit circle (i.e., $r = 1$), then $z^n = (1, n\theta)$.



The n -th Complex Roots of Unity



Solutions to the equation $z^n = 1$

- by the multiplication rules: solutions are $z = (1, \theta)$, for θ a multiple of $2\pi/n$.

The n -th Complex Roots of Unity

Solutions to the equation $z^n = 1$

- by the multiplication rules: solutions are $z = (1, \theta)$, for θ a multiple of $2\pi/n$.
- It can be represented as

$$1, \omega, \omega^2, \dots, \omega^{n-1}$$

where

$$\omega = e^{2\pi i/n}$$



The n -th Complex Roots of Unity

Solutions to the equation $z^n = 1$

- by the multiplication rules: solutions are $z = (1, \theta)$, for θ a multiple of $2\pi/n$.
- It can be represented as


$$1, \omega, \omega^2, \dots, \omega^{n-1}$$

where

$$\omega = e^{2\pi i/n}$$

For n is **even**:

- These numbers are plus-minus paired.
- Their squares are the $(n/2)$ -nd roots of unity.



Complex Conjugate



The **complex conjugate** of a complex number $z = re^{i\theta}$ is $z^* = re^{-i\theta}$.

Complex Conjugate



The **complex conjugate** of a complex number $z = re^{i\theta}$ is $z^* = re^{-i\theta}$.

The **complex conjugate** of a vector (or a matrix) is obtained by taking the complex conjugates of all its entries.

Complex Conjugate



The **complex conjugate** of a complex number $z = re^{i\theta}$ is $z^* = re^{-i\theta}$.

The **complex conjugate** of a vector (or a matrix) is obtained by taking the complex conjugates of all its entries.

The **angle** between two vectors $u = (u_0, \dots, u_{n-1})$ and $v = (v_0, \dots, v_{n-1})$ in $\mathbb{C}^n$ is just a **scaling factor** times their **inner product**

$$u \cdot v^* = u_0 v_0^* + u_1 v_1^* + \dots + u_{n-1} v_{n-1}^*$$

Complex Conjugate



The **complex conjugate** of a complex number $z = re^{i\theta}$ is $z^* = re^{-i\theta}$.

The **complex conjugate** of a vector (or a matrix) is obtained by taking the complex conjugates of all its entries.

The **angle** between two vectors $u = (u_0, \dots, u_{n-1})$ and $v = (v_0, \dots, v_{n-1})$ in $\mathbb{C}^n$ is just a **scaling factor** times their **inner product**

$$u \cdot v^* = u_0 v_0^* + u_1 v_1^* + \dots + u_{n-1} v_{n-1}^*$$

The above quantity is maximized when the vectors lie in the **same direction** and is zero when the vectors are **orthogonal** to each other.

The Fast Fourier Transform

Polynomial multiplication



If $A(x) = a_0 + a_1x + \dots + a_dx^d$ and $B(x) = b_0 + b_1x + \dots + b_dx^d$, their product

$$C(x) = c_0 + c_1x + \dots + c_{2d}x^{2d}$$

has coefficients

$$c_k = a_0b_k + a_1b_{k-1} + \dots + a_kb_0 = \sum_{i=0}^k a_ib_{k-i}$$

where for $i > d$, take a_i and b_i to be zero.

Polynomial multiplication



If $A(x) = a_0 + a_1x + \dots + a_dx^d$ and $B(x) = b_0 + b_1x + \dots + b_dx^d$, their product

$$C(x) = c_0 + c_1x + \dots + c_{2d}x^{2d}$$

has coefficients

$$c_k = a_0b_k + a_1b_{k-1} + \dots + a_kb_0 = \sum_{i=0}^k a_ib_{k-i}$$

where for $i > d$, take a_i and b_i to be zero.

Computing c_k from this formula take $O(k)$ step, and finding all $2d + 1$ coefficients would therefore seem to require $\Theta(d^2)$ time.

Polynomial multiplication



If $A(x) = a_0 + a_1x + \dots + a_dx^d$ and $B(x) = b_0 + b_1x + \dots + b_dx^d$, their product

$$C(x) = c_0 + c_1x + \dots + c_{2d}x^{2d}$$

has coefficients

$$c_k = a_0b_k + a_1b_{k-1} + \dots + a_kb_0 = \sum_{i=0}^k a_ib_{k-i}$$

where for $i > d$, take a_i and b_i to be zero.

Computing c_k from this formula take $O(k)$ step, and finding all $2d + 1$ coefficients would therefore seem to require $\Theta(d^2)$ time.

Q: Can we do better?

An alternative representation

Fact: A degree- d polynomial is uniquely characterized by its values at any $d + 1$ distinct points.

An alternative representation



Fact: A degree- d polynomial is uniquely characterized by its values at any $d + 1$ distinct points.

We can specify a degree- d polynomial $A(x) = a_0 + a_1x + \dots + a_dx^d$ by either of the following:

- Its **coefficients** $a_0, a_1, \dots, a_d$. (coefficient representation).

An alternative representation



Fact: A degree- d polynomial is uniquely characterized by its values at any $d + 1$ distinct points.

We can specify a degree- d polynomial $A(x) = a_0 + a_1x + \dots + a_dx^d$ by either of the following:


- Its **coefficients** $a_0, a_1, \dots, a_d$. (coefficient representation).
- The **values** $A(x_0), A(x_1), \dots, A(x_d)$ (value representation).

An alternative representation


Fact: A degree- d polynomial is uniquely characterized by its values at any $d + 1$ distinct points.

We can specify a degree- d polynomial $A(x) = a_0 + a_1x + \dots + a_dx^d$ by either of the following:

- Its **coefficients** $a_0, a_1, \dots, a_d$. (coefficient representation).
- The **values** $A(x_0), A(x_1), \dots, A(x_d)$ (value representation).




An alternative representation



The product $C(x)$ has degree $2d$, it is determined by its value at any $2d + 1$ points.


An alternative representation



The product $C(x)$ has degree $2d$, it is determined by its value at any $2d + 1$ points.

Its value at any given point z is just $A(z)$ times $B(z)$.

An alternative representation



The product $C(x)$ has degree $2d$, it is determined by its value at any $2d + 1$ points.

Its value at any given point z is just $A(z)$ times $B(z)$.

Therefore, polynomial multiplication takes linear time in the value representation.

The algorithm



Input: Coefficients of two polynomials, $A(x)$ and $B(x)$, of degree d

Output: Their product $C = A \cdot B$

The algorithm



Input: Coefficients of two polynomials, $A(x)$ and $B(x)$, of degree d

Output: Their product $C = A \cdot B$

Selection

Pick some points $x_0, x_1, \dots, x_{n-1}$, where $n \geq 2d + 1$.



The algorithm

Input: Coefficients of two polynomials, $A(x)$ and $B(x)$, of degree d

Output: Their product $C = A \cdot B$

Selection

Pick some points $x_0, x_1, \dots, x_{n-1}$, where $n \geq 2d + 1$.

Evaluation

Compute $A(x_0), A(x_1), \dots, A(x_{n-1})$ and $B(x_0), B(x_1), \dots, B(x_{n-1})$.



The algorithm

Input: Coefficients of two polynomials, $A(x)$ and $B(x)$, of degree d

Output: Their product $C = A \cdot B$

Selection

Pick some points $x_0, x_1, \dots, x_{n-1}$, where $n \geq 2d + 1$.

Evaluation

Compute $A(x_0), A(x_1), \dots, A(x_{n-1})$ and $B(x_0), B(x_1), \dots, B(x_{n-1})$.

Multiplication

Compute $C(x_k) = A(x_k)B(x_k)$ for all $k = 0, \dots, n - 1$.



The algorithm

Input: Coefficients of two polynomials, $A(x)$ and $B(x)$, of degree d

Output: Their product $C = A \cdot B$

Selection

Pick some points $x_0, x_1, \dots, x_{n-1}$, where $n \geq 2d + 1$.

Evaluation

Compute $A(x_0), A(x_1), \dots, A(x_{n-1})$ and $B(x_0), B(x_1), \dots, B(x_{n-1})$.

Multiplication

Compute $C(x_k) = A(x_k)B(x_k)$ for all $k = 0, \dots, n - 1$.

Interpolation

Recover $C(x) = c_0 + c_1x + \dots + c_{2d}x^{2d}$

Fast Fourier Transform



SHANGHAI JIAO TONG
UNIVERSITY

The **selection** step and the **multiplications** are just linear time:

Fast Fourier Transform



The **selection** step and the **multiplications** are just linear time:

- In a **typical setting** for polynomial multiplication, the coefficients of the polynomials are real number.



The **selection** step and the **multiplications** are just linear time:

- In a **typical setting** for polynomial multiplication, the coefficients of the polynomials are real number.
- Moreover, are small enough that basic arithmetic operations take **unit time**.



The **selection** step and the **multiplications** are just linear time:

- In a **typical setting** for polynomial multiplication, the coefficients of the polynomials are real number.
- Moreover, are small enough that basic arithmetic operations take **unit time**.

Evaluating a polynomial of degree $d \leq n$ at a single point takes $O(n)$, and so the baseline for n points is $\Theta(n^2)$.



The **selection** step and the **multiplications** are just linear time:

- In a **typical setting** for polynomial multiplication, the coefficients of the polynomials are real number.
- Moreover, are small enough that basic arithmetic operations take **unit time**.

Evaluating a polynomial of degree $d \leq n$ at a single point takes $O(n)$, and so the baseline for n points is $\Theta(n^2)$.

The **Fast Fourier Transform (FFT)** does it in just $O(n \log n)$ time, for a particularly clever choice of $x_0, \dots, x_{n-1}$.

Evaluation by divide-and-conquer



SHANGHAI JIAO TONG
UNIVERSITY

Q: How to make it efficient?

Evaluation by divide-and-conquer



Q: How to make it efficient?

First idea, we pick the n points,

$$\pm x_0, \pm x_1, \dots, \pm x_{n/2-1}$$

then the computations required for each $A(x_i)$ and $A(-x_i)$ overlap a lot, because the **even power** of x_i coincide with those of $-x_i$.

Evaluation by divide-and-conquer



Q: How to make it efficient?

First idea, we pick the n points,

$$\pm x_0, \pm x_1, \dots, \pm x_{n/2-1}$$

then the computations required for each $A(x_i)$ and $A(-x_i)$ overlap a lot, because the **even power** of x_i coincide with those of $-x_i$.

We need to split $A(x)$ into its odd and even powers, for instance

$$3 + 4x + 6x^2 + 2x^3 + x^4 + 10x^5 = (3 + 6x^2 + x^4) + x(4 + 2x^2 + 10x^4)$$

Evaluation by divide-and-conquer



Q: How to make it efficient?

First idea, we pick the n points,

$$\pm x_0, \pm x_1, \dots, \pm x_{n/2-1}$$

then the computations required for each $A(x_i)$ and $A(-x_i)$ overlap a lot, because the **even power** of x_i coincide with those of $-x_i$.

We need to split $A(x)$ into its odd and even powers, for instance

$$3 + 4x + 6x^2 + 2x^3 + x^4 + 10x^5 = (3 + 6x^2 + x^4) + x(4 + 2x^2 + 10x^4)$$

More generally

$$A(x) = A_e(x^2) + xA_o(x^2)$$

where $A_e(\cdot)$, with the **even-numbered coefficients**, and $A_o(\cdot)$, with the **odd-numbered coefficients**, are polynomials of degree $\leq n/2 - 1$.

Evaluation by divide-and-conquer



Given paired points $\pm x_i$, the calculations needed for $A(x_i)$ can be recycled toward computing $A(-x_i)$:

$$A(x_i) = A_e(x_i^2) + x_i A_o(x_i^2)$$

$$A(-x_i) = A_e(x_i^2) - x_i A_o(x_i^2)$$

Evaluation by divide-and-conquer



Given paired points $\pm x_i$, the calculations needed for $A(x_i)$ can be recycled toward computing $A(-x_i)$:

$$A(x_i) = A_e(x_i^2) + x_i A_o(x_i^2)$$

$$A(-x_i) = A_e(x_i^2) - x_i A_o(x_i^2)$$

Evaluating $A(x)$ at n paired points $\pm x_0, \dots, \pm x_{n/2-1}$ reduces to evaluating $A_e(x)$ and $A_o(x)$ at just $n/2$ points, $x_0^2, \dots, x_{n/2-1}^2$.

Evaluation by divide-and-conquer



Given paired points $\pm x_i$, the calculations needed for $A(x_i)$ can be recycled toward computing $A(-x_i)$:

$$A(x_i) = A_e(x_i^2) + x_i A_o(x_i^2)$$

$$A(-x_i) = A_e(x_i^2) - x_i A_o(x_i^2)$$

Evaluating $A(x)$ at n paired points $\pm x_0, \dots, \pm x_{n/2-1}$ reduces to evaluating $A_e(x)$ and $A_o(x)$ at just $n/2$ points, $x_0^2, \dots, x_{n/2-1}^2$.

If we could recurse, we would get a divide-and-conquer procedure with running time

$$T(n) = 2T(n/2) + O(n) = O(n \log n)$$

How to choose n points?



Aim: To recurse at the next level, we need the $n/2$ evaluation points $x_0^2, x_1^2, \dots, x_{n/2-1}^2$ to be themselves **plus-minus pairs**.

How to choose n points?



Aim: To recurse at the next level, we need the $n/2$ evaluation points $x_0^2, x_1^2, \dots, x_{n/2-1}^2$ to be themselves **plus-minus pairs**.

Q: How can a square be negative?

How to choose n points?



Aim: To recurse at the next level, we need the $n/2$ evaluation points $x_0^2, x_1^2, \dots, x_{n/2-1}^2$ to be themselves **plus-minus pairs**.

Q: How can a square be negative?

- We use **complex numbers**.



How to choose n points?

Aim: To recurse at the next level, we need the $n/2$ evaluation points $x_0^2, x_1^2, \dots, x_{n/2-1}^2$ to be themselves **plus-minus pairs**.

Q: How can a square be negative?

- We use **complex numbers**.

At the very bottom of the recursion, we have a **single point**, 1 , in which case the level above it must consist of its square roots, $\pm\sqrt{1} = \pm 1$.



How to choose n points?

Aim: To recurse at the next level, we need the $n/2$ evaluation points $x_0^2, x_1^2, \dots, x_{n/2-1}^2$ to be themselves **plus-minus pairs**.

Q: How can a square be negative?

- We use **complex numbers**.

At the very bottom of the recursion, we have a **single point**, 1 , in which case the level above it must consist of its square roots, $\pm\sqrt{1} = \pm 1$.

The next level up then has $\pm\sqrt{+1} = \pm 1$, as well as the complex numbers $\pm\sqrt{-1} = \pm i$.



How to choose n points?

Aim: To recurse at the next level, we need the $n/2$ evaluation points $x_0^2, x_1^2, \dots, x_{n/2-1}^2$ to be themselves **plus-minus pairs**.

Q: How can a square be negative?

- We use **complex numbers**.

At the very bottom of the recursion, we have a **single point**, 1 , in which case the level above it must consist of its square roots, $\pm\sqrt{1} = \pm 1$.

The next level up then has $\pm\sqrt{+1} = \pm 1$, as well as the complex numbers $\pm\sqrt{-1} = \pm i$.

By continuing in this manner, we eventually reach the initial set of n points: the **complex n th** roots of unity, that is the n complex solutions of the equation

$$z^n = 1$$

SHANGHAI JIAO TONG
UNIVERSITY

- by the multiplication rules: solutions are $z = (1, \theta)$, for θ a multiple of $2\pi/n$.
- It can be represented as

$$1, \omega, \omega^2, \dots, \omega^{n-1}$$

$$\omega = e^{2\pi i/n}$$



The n -th complex roots of unity

Solutions to the equation $z^n = 1$

- by the multiplication rules: solutions are $z = (1, \theta)$, for θ a multiple of $2\pi/n$.
- It can be represented as

$$1, \omega, \omega^2, \dots, \omega^{n-1}$$

where

$$\omega = e^{2\pi i/n}$$

For n is even:

- These numbers are plus-minus paired.
- Their squares are the $(n/2)$ -nd roots of unity.



The FFT algorithm

FFT (A, ω)

input : coefficient representation of a polynomial $A(x)$ of degree $\leq n - 1$, where n is a power of 2;
 ω , an n -th root of unity

output: value representation $A(\omega^0), \dots, A(\omega^{n-1})$

if $\omega = 1$ **then** return $A(1)$;

express $A(x)$ in the form $A_e(x^2) + xA_o(x^2)$;

call FFT (A_e, ω^2) to evaluate A_e at even powers of ω ;

call FFT (A_o, ω^2) to evaluate A_o at even powers of ω ;

for $j = 0$ **to** $n - 1$ **do**

 | compute $A(\omega^j) = A_e(\omega^{2j}) + \omega^j A_o(\omega^{2j})$;

end

return ($A(\omega^0), \dots, A(\omega^{n-1})$);

Interpolation



FFT moves from **coefficients** to **values** in time just $O(n \log n)$, when the points $\{x_i\}$ are complex n -th roots of unity $(1, \omega, \omega^2, \dots, \omega^{n-1})$.

Interpolation



FFT moves from **coefficients** to **values** in time just $O(n \log n)$, when the points $\{x_i\}$ are complex n -th roots of unity $(1, \omega, \omega^2, \dots, \omega^{n-1})$.

That is,

$$\langle value \rangle = \text{FFT}(\langle coefficients \rangle, \omega)$$



FFT moves from **coefficients** to **values** in time just $O(n \log n)$, when the points $\{x_i\}$ are complex n -th roots of unity $(1, \omega, \omega^2, \dots, \omega^{n-1})$.

That is,

$$\langle value \rangle = \text{FFT}(\langle coefficients \rangle, \omega)$$

We will see that the interpolation can be computed by

$$\langle coefficients \rangle = \frac{1}{n} \text{FFT}(\langle values \rangle, \omega^{-1})$$

A matrix reformation



Let's explicitly set down the relationship between our two representations for a polynomial $A(x)$ of degree $\leq n - 1$.

A matrix reformation



Let's explicitly set down the relationship between our two representations for a polynomial $A(x)$ of degree $\leq n - 1$.

$$\begin{bmatrix} A(x_0) \\ A(x_1) \\ \vdots \\ A(x_{n-1}) \end{bmatrix} = \begin{bmatrix} 1 & x_0 & x_0^2 & \dots & x_0^{n-1} \\ 1 & x_1 & x_1^2 & \dots & x_1^{n-1} \\ & & \vdots & & \\ 1 & x_{n-1} & x_{n-1}^2 & \dots & x_{n-1}^{n-1} \end{bmatrix} \begin{bmatrix} a_0 \\ a_1 \\ \vdots \\ a_{n-1} \end{bmatrix}$$

A matrix reformation



Let's explicitly set down the relationship between our two representations for a polynomial $A(x)$ of degree $\leq n - 1$.

$$\begin{bmatrix} A(x_0) \\ A(x_1) \\ \vdots \\ A(x_{n-1}) \end{bmatrix} = \begin{bmatrix} 1 & x_0 & x_0^2 & \dots & x_0^{n-1} \\ 1 & x_1 & x_1^2 & \dots & x_1^{n-1} \\ & & \vdots & & \\ 1 & x_{n-1} & x_{n-1}^2 & \dots & x_{n-1}^{n-1} \end{bmatrix} \begin{bmatrix} a_0 \\ a_1 \\ \vdots \\ a_{n-1} \end{bmatrix}$$

Let M be the matrix in the middle, which is a **Vandermonde** matrix.

A matrix reformation



Let's explicitly set down the relationship between our two representations for a polynomial $A(x)$ of degree $\leq n-1$.

$$\begin{bmatrix} A(x_0) \\ A(x_1) \\ \vdots \\ A(x_{n-1}) \end{bmatrix} = \begin{bmatrix} 1 & x_0 & x_0^2 & \dots & x_0^{n-1} \\ 1 & x_1 & x_1^2 & \dots & x_1^{n-1} \\ & & \vdots & & \\ 1 & x_{n-1} & x_{n-1}^2 & \dots & x_{n-1}^{n-1} \end{bmatrix} \begin{bmatrix} a_0 \\ a_1 \\ \vdots \\ a_{n-1} \end{bmatrix}$$

Let M be the matrix in the middle, which is a **Vandermonde** matrix.

- If $x_0, x_1, \dots, x_{n-1}$ are distinct numbers, then M is invertible.

A matrix reformation



Let's explicitly set down the relationship between our two representations for a polynomial $A(x)$ of degree $\leq n-1$.

$$\begin{bmatrix} A(x_0) \\ A(x_1) \\ \vdots \\ A(x_{n-1}) \end{bmatrix} = \begin{bmatrix} 1 & x_0 & x_0^2 & \dots & x_0^{n-1} \\ 1 & x_1 & x_1^2 & \dots & x_1^{n-1} \\ & & \vdots & & \\ 1 & x_{n-1} & x_{n-1}^2 & \dots & x_{n-1}^{n-1} \end{bmatrix} \begin{bmatrix} a_0 \\ a_1 \\ \vdots \\ a_{n-1} \end{bmatrix}$$

Let M be the matrix in the middle, which is a **Vandermonde** matrix.

- If $x_0, x_1, \dots, x_{n-1}$ are distinct numbers, then M is invertible.
- **evaluation** is multiplication by M , while **interpolation** is multiplication by M^{-1} .

A matrix reformation



This **reformulation** of our polynomial operations reveals their essential nature more clearly.

A matrix reformation



This **reformulation** of our polynomial operations reveals their essential nature more clearly.

It justifies an assumption that $A(x)$ is **uniquely characterized** by its values at any n points.

A matrix reformation



This **reformulation** of our polynomial operations reveals their essential nature more clearly.

It justifies an assumption that $A(x)$ is **uniquely characterized** by its values at any n points.

Vandermonde matrices also have the distinction of being quicker to invert than more general matrices, in $O(n^2)$ time instead of $O(n^3)$.

A matrix reformation



This **reformulation** of our polynomial operations reveals their essential nature more clearly.

It justifies an assumption that $A(x)$ is **uniquely characterized** by its values at any n points.

Vandermonde matrices also have the distinction of being quicker to invert than more general matrices, in $O(n^2)$ time instead of $O(n^3)$.

However, using this for interpolation would still not be fast enough for us..

Interpolation resolved



In linear algebra terms, the FFT multiplies an arbitrary n -dimensional vector, which we have been calling the coefficient representation, by the $n \times n$ matrix.

$$M_n(\omega) = \begin{bmatrix} 1 & 1 & 1 & \dots & 1 \\ 1 & \omega & \omega^2 & \dots & \omega^{n-1} \\ & & \vdots & & \\ 1 & \omega^j & \omega^{2j} & \dots & \omega^{(n-1)j} \\ & & \vdots & & \\ 1 & \omega^{n-1} & \omega^{2(n-1)} & \dots & \omega^{(n-1)(n-1)} \end{bmatrix}$$

Its (j, k) -th entry (starting row- and column-count at zero) is ω^{jk}

Interpolation resolved



The columns of M are orthogonal to each other, which is often called the Fourier basis.

Interpolation resolved



The columns of M are orthogonal to each other, which is often called the Fourier basis.

The FFT is thus a change of basis, a rigid rotation. The inverse of M is the opposite rotation, from the Fourier basis back into the standard basis.

Interpolation resolved



The columns of M are orthogonal to each other, which is often called the Fourier basis.

The FFT is thus a change of basis, a rigid rotation. The inverse of M is the opposite rotation, from the Fourier basis back into the standard basis.

Inversion formula

$$M_n(\omega)^{-1} = \frac{1}{n} M_n(\omega^{-1})$$

Interpolation resolved



Take ω to be $e^{2\pi i/n}$, and think of M as vectors in $\mathbb{C}^n$.

Interpolation resolved



Take ω to be $e^{2\pi i/n}$, and think of M as **vectors** in $\mathbb{C}^n$.

Recall that the **angle** between two vectors $u = (u_0, \dots, u_{n-1})$ and $v = (v_0, \dots, v_{n-1})$ in $\mathbb{C}^n$ is just a **scaling factor** times their **inner product**

$$u \cdot v^* = u_0 v_0^* + u_1 v_1^* + \dots + u_{n-1} v_{n-1}^*$$

where z^* denotes the **complex conjugate** of z .

Interpolation resolved



Take ω to be $e^{2\pi i/n}$, and think of M as **vectors** in $\mathbb{C}^n$.

Recall that the **angle** between two vectors $u = (u_0, \dots, u_{n-1})$ and $v = (v_0, \dots, v_{n-1})$ in $\mathbb{C}^n$ is just a **scaling factor** times their **inner product**

$$u \cdot v^* = u_0 v_0^* + u_1 v_1^* + \dots + u_{n-1} v_{n-1}^*$$

where z^* denotes the **complex conjugate** of z .

The above quantity is maximized when the vectors lie in the **same direction** and is zero when the vectors are **orthogonal** to each other.

Interpolation resolved



Lemma

The columns of matrix M are orthogonal to each other.

Interpolation resolved



Lemma

The columns of matrix M are orthogonal to each other.

Proof.



Lemma

The columns of matrix M are orthogonal to each other.

Proof.

- Take the inner product of any columns j and k of matrix M ,

$$1 + \omega^{j-k} + \omega^{2(j-k)} + \dots + \omega^{(n-1)(j-k)}$$

This is a **geometric series** with first term 1, last term $\omega^{(n-1)(j-k)}$, and ratio ω^{j-k} .



Lemma

The columns of matrix M are orthogonal to each other.

Proof.

- Take the inner product of any columns j and k of matrix M ,

$$1 + \omega^{j-k} + \omega^{2(j-k)} + \dots + \omega^{(n-1)(j-k)}$$

This is a **geometric series** with first term 1, last term $\omega^{(n-1)(j-k)}$, and ratio ω^{j-k} .

- Therefore, if $j \neq k$, it evaluates to

$$\frac{1 - \omega^{n(j-k)}}{1 - \omega^{(j-k)}} = 0$$



Lemma

The columns of matrix M are orthogonal to each other.

Proof.

- Take the inner product of any columns j and k of matrix M ,

$$1 + \omega^{j-k} + \omega^{2(j-k)} + \dots + \omega^{(n-1)(j-k)}$$

This is a **geometric series** with first term 1, last term $\omega^{(n-1)(j-k)}$, and ratio ω^{j-k} .

- Therefore, if $j \neq k$, it evaluates to

$$\frac{1 - \omega^{n(j-k)}}{1 - \omega^{(j-k)}} = 0$$

- If $j = k$, then it evaluates to n .



Corollary

$MM^* = nI$, i.e.,

$$M_n^{-1} = \frac{1}{n} M_n^*$$

The definitive FFT algorithm



The FFT takes as input a **vector** $a = (a_0, \dots, a_{n-1})$ and a **complex number** ω whose powers $1, \omega, \omega^2, \dots, \omega^{n-1}$ are the complex n -th roots of unity.

The definitive FFT algorithm



The FFT takes as input a **vector** $a = (a_0, \dots, a_{n-1})$ and a **complex number** ω whose powers $1, \omega, \omega^2, \dots, \omega^{n-1}$ are the complex n -th roots of unity.

It **multiplies** vector a by the $n \times n$ matrix $M_n(\omega)$, which has (j, k) -th entry ω^{jk} .

The definitive FFT algorithm



The FFT takes as input a **vector** $a = (a_0, \dots, a_{n-1})$ and a **complex number** ω whose powers $1, \omega, \omega^2, \dots, \omega^{n-1}$ are the complex n -th roots of unity.

It **multiplies** vector a by the $n \times n$ matrix $M_n(\omega)$, which has (j, k) -th entry ω^{jk} .

The potential for using divide-and-conquer in this matrix-vector multiplication becomes apparent when M 's columns are segregated into evens and odds.

The definitive FFT algorithm



The FFT takes as input a **vector** $a = (a_0, \dots, a_{n-1})$ and a **complex number** ω whose powers $1, \omega, \omega^2, \dots, \omega^{n-1}$ are the complex n -th roots of unity.

It **multiplies** vector a by the $n \times n$ matrix $M_n(\omega)$, which has (j, k) -th entry ω^{jk} .

The potential for using divide-and-conquer in this matrix-vector multiplication becomes apparent when M 's columns are segregated into evens and odds.

The product of $M_n(\omega)$ with vector $a = (a_0, \dots, a_{n-1})$, a size- n problem, can be expressed in terms of two size- $n/2$ problems: the product of $M_{n/2}(\omega^2)$ with $(a_0, a_2, \dots, a_{n-2})$ and with $(a_1, a_3, \dots, a_{n-1})$.



The definitive FFT algorithm

The FFT takes as input a **vector** $a = (a_0, \dots, a_{n-1})$ and a **complex number** ω whose powers $1, \omega, \omega^2, \dots, \omega^{n-1}$ are the complex n -th roots of unity.

It **multiplies** vector a by the $n \times n$ matrix $M_n(\omega)$, which has (j, k) -th entry ω^{jk} .

The potential for using divide-and-conquer in this matrix-vector multiplication becomes apparent when M 's columns are segregated into evens and odds.

The product of $M_n(\omega)$ with vector $a = (a_0, \dots, a_{n-1})$, a size- n problem, can be expressed in terms of two size- $n/2$ problems: the product of $M_{n/2}(\omega^2)$ with $(a_0, a_2, \dots, a_{n-2})$ and with $(a_1, a_3, \dots, a_{n-1})$.

This divide-and-conquer strategy leads to the definitive FFT algorithm, whose running time is $T(n) = 2T(n/2) + O(n) = O(n \log n)$.

The general FFT algorithm



FFT (a, ω)

input : An array $a = (a_0, a_1, \dots, a_{n-1})$ for n is a power of 2; ω , an n -th root of unity

output: $M_n(\omega)a$

if $\omega = 1$ **then** return a ;

$(s_0, s_1, \dots, s_{n/2-1}) = \text{FFT}((a_0, a_2, \dots, a_{n-2}), \omega^2)$;

$(s'_0, s'_1, \dots, s'_{n/2-1}) = \text{FFT}((a_1, a_3, \dots, a_{n-1}), \omega^2)$;

for $j = 0$ **to** $n/2 - 1$ **do**

$r_j = s_j + \omega^j s'_j$;
 $r_{j+n/2} = s_j - \omega^j s'_j$;

end

return $(r_0, r_1, \dots, r_{n-1})$;

Top 10 algorithms of the 20th century



SHANGHAI JIAO TONG
UNIVERSITY

Top 10 algorithms of the 20th century



1946: The Metropolis Algorithm

1947: Simplex Method

1950: Krylov Subspace Method

1951: The Decompositional Approach to Matrix Computations

1957: The Fortran Optimizing Compiler

1959: QR Algorithm

1962: Quicksort

1965: Fast Fourier Transform

1977: Integer Relation Detection

1987: Fast Multipole Method

Top 10 algorithms of the 20th century



SHANGHAI JIAO TONG
UNIVERSITY

Top 10 algorithms of the 20th century



SHANGHAI JIAO TONG
UNIVERSITY

1946: The Metropolis Algorithm

1947: Simplex Method

1950: Krylov Subspace Method

1951: The Decompositional Approach to Matrix Computations

1957: The Fortran Optimizing Compiler

1959: QR Algorithm

1962: Quicksort

1965: Fast Fourier Transform

1977: Integer Relation Detection

1987: Fast Multipole Method